

[Become a Host](#)

[Help](#)

[Sign Up](#)

[Log In](#)

# LIVE THERE

Book homes from local hosts in 191+ countries and experience a place like you live there.



Presented by: Pranati Sahu & Subu Govindaswamy



## Airbnb New User Bookings

Wed 25 Nov 2015 – Thu 11 Feb 2016 (3 months ago)

Airbnb challenges to predict in which country a new user will make his or her first booking.

# Data Sources:

All “United States” user data from 2010 - 2014.

- ★ train\_users.csv
- ★ sessions.csv
- ★ countries.csv
- ★ Age\_gender\_bkts.csv

Source: <https://www.kaggle.com/c/airbnb-recruiting-new-user-bookings>

# Data Samples

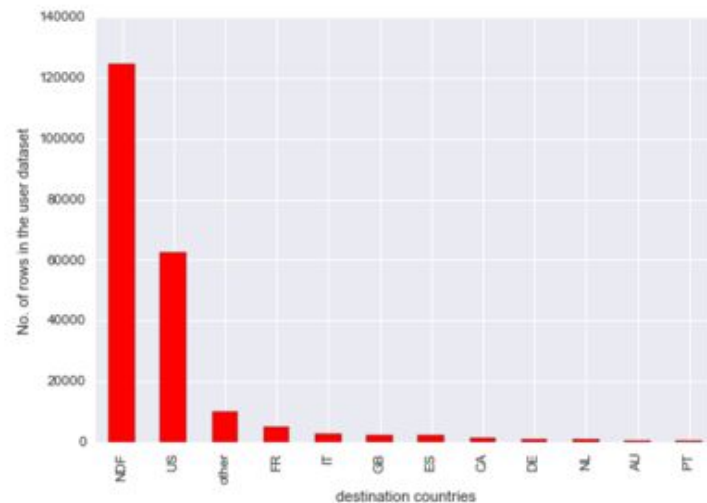
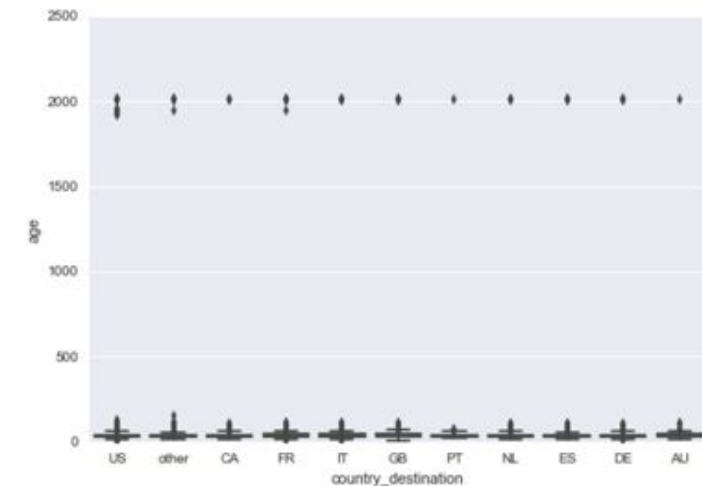
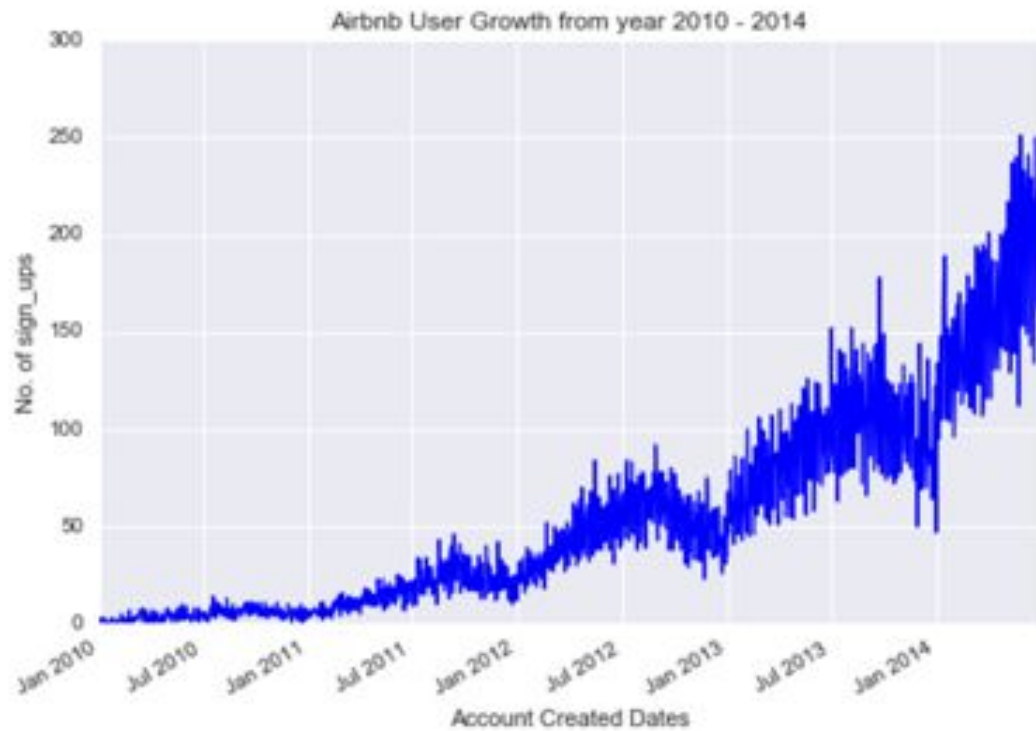
**Ugly Slide Alert!!!**

	id	date_account_created	timestamp_first_active	date_first_booking	gender	age	signup_method	signup_flow	language	affiliat
0	gxn3p5htnn	2010-06-28	20090319043255	NaN	- unknown-	NaN	facebook	0	en	direct
1	820tgsjq7	2011-05-25	20090523174809	NaN	MALE	38.0	facebook	0	en	seo

	country_destination	lat_destination	lng_destination	distance_km	destination_km2	destination_language	language_levenshtein_distance
0	AU	-26.853388	133.275160	15297.7440	7741220.0	eng	0.00
1	CA	62.393303	-96.818146	2828.1333	9984670.0	eng	0.00

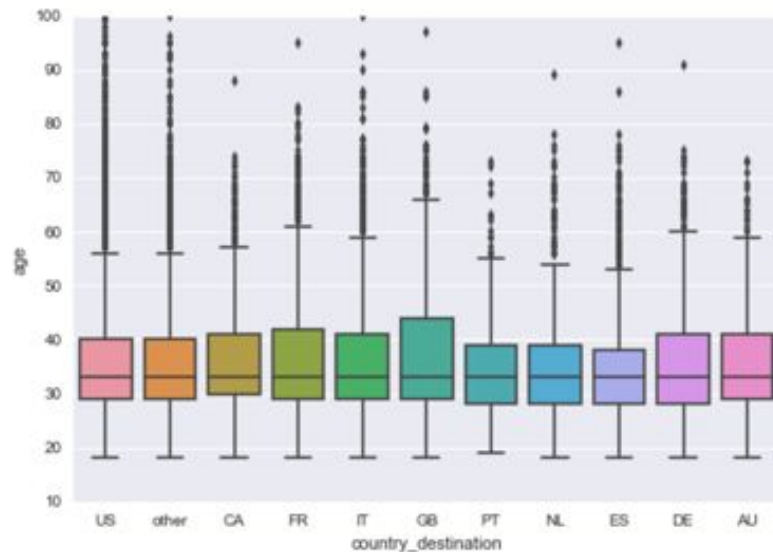
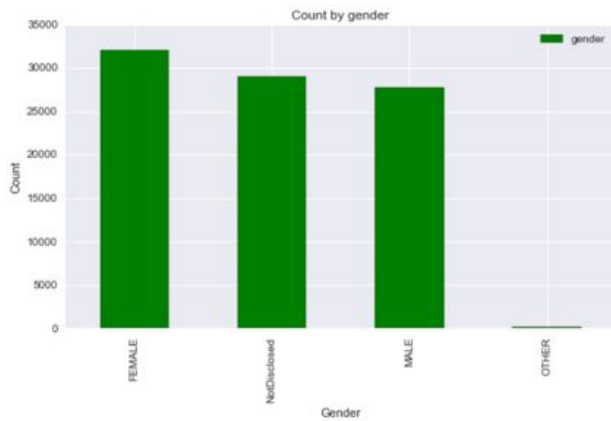
	user_id	action	action_type	action_detail	device_type	secs_elapsed
0	d1mm9tcy42	lookup	NaN	NaN	Windows Desktop	319.0
1	d1mm9tcy42	search_results	click	view_search_results	Windows Desktop	67753.0
2	d1mm9tcy42	lookup	NaN	NaN	Windows Desktop	301.0

# Exploration Data Analysis



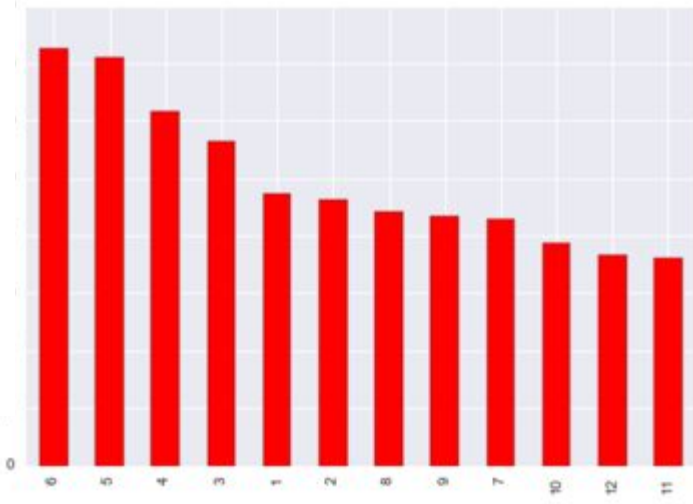
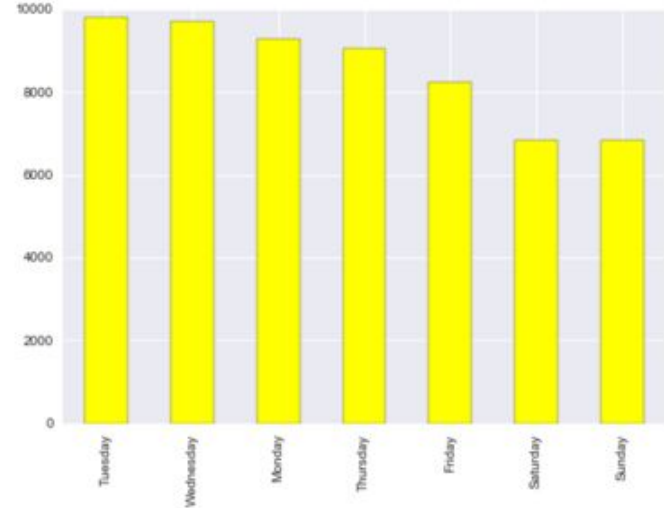
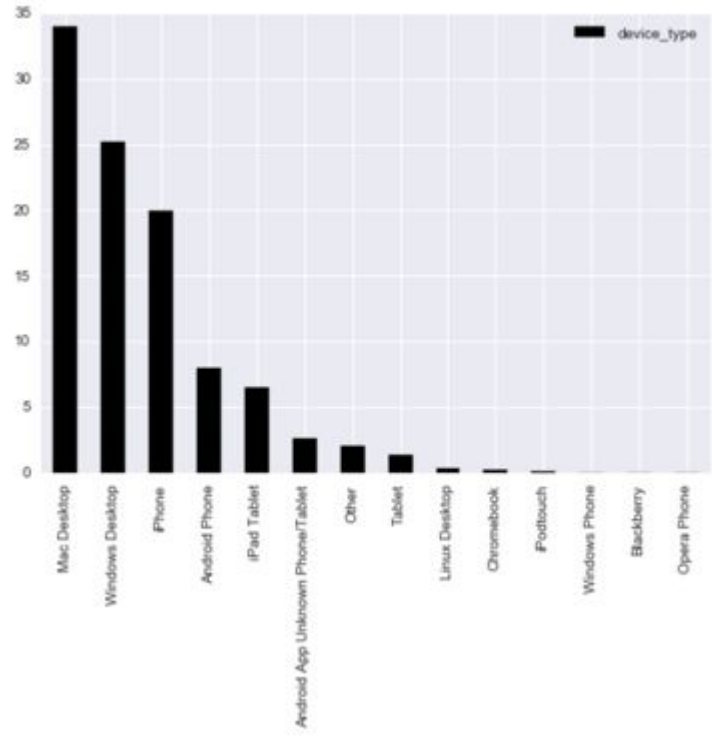
# Feature Engineering

- ★ ~32% users did not disclose their gender
- ★ ~30% users did not declare their age



```
In [45]: # Make a new year/month/day column for date first booking
users['dfb_year'] = users['date_first_booking'].apply(lambda x:x.year)
users['dfb_month'] = users['date_first_booking'].apply(lambda x:x.month)
users['dfb_day'] = users['date_first_booking'].apply(lambda x:x.day)
users['dfb_wn'] = users['date_first_booking'].apply(lambda x:x.isocalendar()[1])
users['dfb_wd'] = users['date_first_booking'].apply(lambda x:calendar.day_name[x.weekday()])
```

# Fun Facts





# Data Munging

	numcountrdesti	age	signup_flow	Numericgenders	Numersignupmeth	Numuserlangu	Numaffchanne	Numaffprovi	numfirststaffrack	numsignupapp	numfirstdevice	numfirstbrowser
numcountrdesti	1.000000	-0.000744	-0.027955	-0.011209	0.118112	-0.022233	-0.049151	-0.044953	-0.056567	-0.041798	-0.042826	-0.036044
age	-0.000744	1.000000	-0.018821	0.016279	0.060980	-0.009287	-0.000122	-0.000752	0.019883	-0.017294	0.008191	0.019352
signup_flow	-0.027955	-0.018821	1.000000	-0.027329	-0.010516	0.002114	-0.012676	-0.008430	-0.138517	0.750494	0.340099	0.267431
Numericgenders	-0.011209	0.016279	-0.027329	1.000000	0.046583	-0.004360	0.019591	0.012176	0.030261	-0.034624	-0.013382	0.059893
Numersignupmeth	0.118112	0.060980	-0.010516	0.046583	1.000000	-0.054856	-0.106041	-0.078732	-0.027920	-0.066080	-0.029534	0.021911
Numuserlangu	-0.022233	-0.009287	0.002114	-0.004360	-0.054856	1.000000	0.050359	0.043598	0.021415	0.006762	0.018893	-0.003194
Numaffchanne	-0.049151	-0.000122	-0.012676	0.019591	-0.106041	0.050359	1.000000	0.662893	0.264107	0.073464	0.047573	0.008791
Numaffprovid	-0.044953	-0.000752	-0.008430	0.012176	-0.078732	0.043598	0.662893	1.000000	0.252985	0.059056	0.041928	0.042434
numfirststaffrack	-0.056567	0.019883	-0.138517	0.030261	-0.027920	0.021415	0.264107	0.252985	1.000000	-0.130181	0.105574	0.038317
numsignupapp	-0.041798	-0.017294	0.750494	-0.034624	-0.066080	0.006762	0.073464	0.059056	-0.130181	1.000000	0.376181	0.250118
numfirstdevice	-0.042826	0.008191	0.340099	-0.013382	-0.029534	0.018893	0.047573	0.041928	0.105574	0.376181	1.000000	0.632981
numfirstbrowser	-0.036044	0.019352	0.267431	0.059893	0.021911	-0.003194	0.008791	0.042434	0.038317	0.250118	0.632981	1.000000



# Correlation Matrix

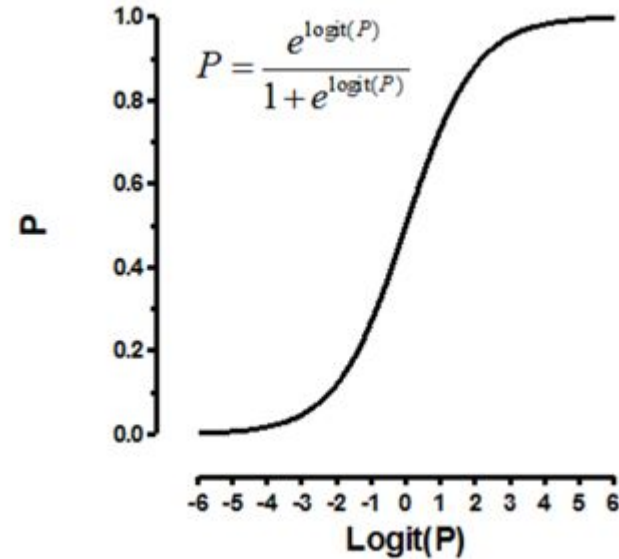
	numcountrede	age	signup_flow	Numericgenders	Numersignupmeth	Numuserlangu	Numaffchanne	Numaffprovid	numfirststafftrack	numsignupapp	numfirstdevice
numcountrede	★ 1.000000										
age	★ -0.000744	★ 1.000000									
signup_flow	★ -0.027955	★ -0.018821	★ 1.000000								
Numericgenders	★ -0.011209	★ 0.016279	★ -0.027329	★ 1.000000							
Numersignupmeth	★ 0.118112	★ 0.060980	★ -0.010516	★ 0.046583	★ 1.000000						
Numuserlangu	★ -0.022233	★ -0.009287	★ 0.002114	★ -0.004360	★ -0.054856	★ 1.000000					
Numaffchanne	★ -0.049151	★ -0.000122	★ -0.012676	★ 0.019591	★ -0.106041	★ 0.050359	★ 1.000000				
Numaffprovid	★ -0.044953	★ -0.000752	★ -0.008430	★ 0.012176	★ -0.078732	★ 0.043598	★ 0.662893	★ 1.000000			
numfirststafftrack	★ -0.056567	★ 0.019883	★ -0.138517	★ 0.030261	★ -0.027920	★ 0.021415	★ 0.264107	★ 0.252985	★ 1.000000		
numsignupapp	★ -0.041798	★ -0.017294	★ 0.750494	★ -0.034624	★ -0.066080	★ 0.006762	★ 0.073464	★ 0.059056	★ -0.130181	★ 1.000000	
numfirstdevice	★ -0.042826	★ 0.008191	★ 0.340099	★ -0.013382	★ -0.029534	★ 0.018893	★ 0.047573	★ 0.041928	★ 0.105574	★ 0.376181	★ 1.000000
numfirstbrowser	★ -0.036044	★ 0.019352	★ 0.267431	★ 0.059893	★ 0.021911	★ -0.003194	★ 0.008791	★ 0.042434	★ 0.038317	★ 0.250118	★ 0.632981

# Findings From Correlations

- ★ Highly correlated variables:
  - Signup\_App and Signup\_Flow (0.75)
  - First\_Browser and First\_Device
- ★ Moderately correlated variables:
  - First\_Browser and Signup\_App
  - First\_Device and Signup\_Flow (0.34)
  - First\_Device and Signup\_App
- ★ No correlations found
  - Predictors and Response variable (country\_destination) (~ 0.00)

# Machine Learning Model Used

Logistics Regression



# Logistic Regression Results

```
In [184]: logreg = LogisticRegression(C=1e9)
feature_cols = ['age', 'Numericgenders', 'signup_flow', 'Numersignupmeth', 'Numaffchanne', 'Numuse
rlangu', 'Numaffprovid']
X = ouserdata_NEW[feature_cols]
y = ouserdata_NEW.numcountrdesti
model = logreg.fit(X, y)
model.score(X, y)
```

```
Out[184]: 0.70157916048049673
```

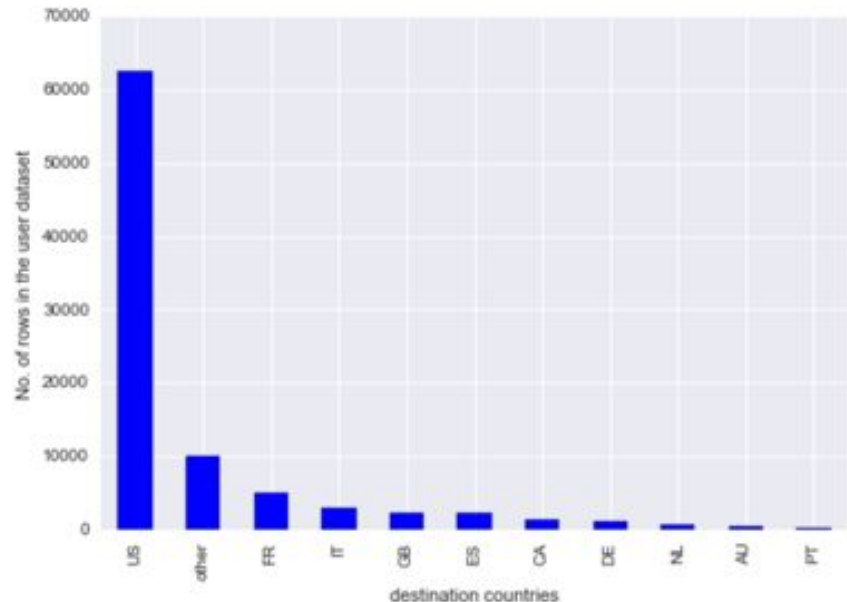
# Null Accuracy Rate

The most common country\_destination is US.

The null accuracy rate = # of rows with USA divided by total number of rows

$$= 62376/88908$$

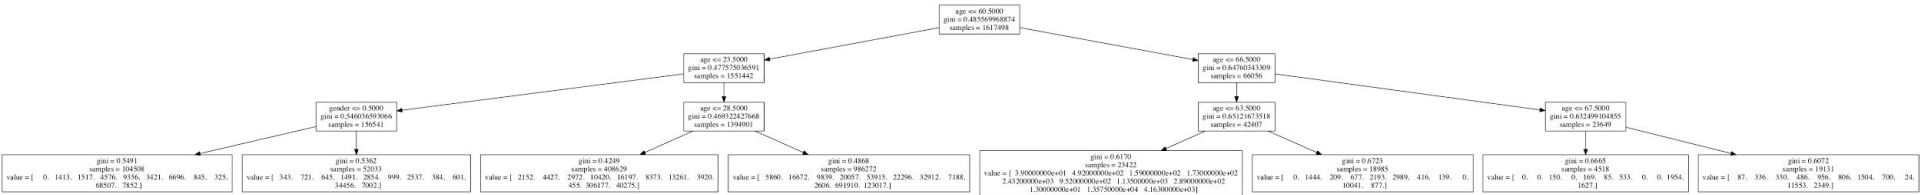
$$= 70.15\%$$



# Machine Learning Models Used

Decision Tree

Random Forest



# Decision Tree Results

```
In [37]: new_df.sort_values(ascending=False, by='importance').head(20)
```

Out[37]:

	feature	importance
0	age	0.143020
461	device_type_Mac Desktop	0.023034
394	signup_app_Web	0.018041
312	dfb_wd_Wednesday	0.017748
341	language_en	0.017284

## MAX TREE DEPTH - 10

```
In [28]: # fit a classification tree with max_depth=10 on all data
from sklearn.tree import DecisionTreeClassifier
treeclf = DecisionTreeClassifier(max_depth=10)
treeclf.fit(X_train, y_train)
```

```
In [32]: # Calculate Accuracy:
from sklearn import metrics
metrics.accuracy_score(y_test, preds)
```

Out[32]: 0.70405862639542494



# Fine Tuned Tree Results

## MAX TREE DEPTH - 100

```
In [6]: # fit a classification tree with max_depth=100 on all data
from sklearn.tree import DecisionTreeClassifier
treeclf10 = DecisionTreeClassifier(max_depth=100)
treeclf10.fit(X_train, y_train)
```

```
Out[6]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=100,
                               max_features=None, max_leaf_nodes=None, min_samples_leaf=1,
                               min_samples_split=2, min_weight_fraction_leaf=0.0,
                               random_state=None, splitter='best')
```

```
In [11]: # Calculate Accuracy:
from sklearn import metrics
metrics.accuracy_score(y_test, preds10)
```

```
Out[11]: 0.97947463318184314
```

	feature	importance
0	age	0.114985
445	secs_elapsed	0.034054
411	first_browser_Chrome	0.023988
391	first_affiliate_tracked_untracked	0.023972
434	first_browser_Safari	0.022352
416	first_browser_Firefox	0.021238
385	first_affiliate_tracked_linked	0.020587
364	affiliate_channel_sem-brand	0.016690
388	first_affiliate_tracked_omg	0.015087
365	affiliate_channel_sem-non-brand	0.014892
314	gender_MALE	0.014197
317	signup_method_basic	0.013681
318	signup_method_facebook	0.013644
315	gender_NotDisclosed	0.013616
313	gender_FEMALE	0.011983
418	first_browser_IE	0.011955
375	affiliate_provider_google	0.009908
402	first_device_type_Windows Desktop	0.009634
366	affiliate_channel_seo	0.009173
371	affiliate_provider_direct	0.008873

# Confusion Matrix

Predicted ->	AU	CA	DE	ES	FR	GB	IT	NL	PT	US	other
Actuals: <b>AU</b>	2937	1	0	3	19	0	2	0	0	83	43
<b>CA</b>	6	9566	11	19	29	7	22	0	1	174	73
<b>DE</b>	0	7	5702	0	24	2	7	0	0	130	51
<b>ES</b>	6	21	0	16111	18	19	27	9	0	376	79
<b>FR</b>	17	18	19	27	35296	11	149	13	0	770	287
<b>GB</b>	0	8	2	16	5	15459	36	4	0	356	96
<b>IT</b>	6	17	6	26	158	28	24441	37	11	692	142
<b>NL</b>	0	2	0	5	15	4	38	5292	0	196	16
<b>PT</b>	0	1	0	0	0	13	12	4	1502	34	16
<b>US</b>	69	188	122	442	734	394	806	194	42	413653	2150
<b>other</b>	31	53	40	64	290	124	188	39	6	2169	77376

# Random Forest Results

```
In [14]: # compute the out-of-bag R-squared score  
rfreg.oob_score_
```

```
Out[14]: 0.93110482776278691
```

```
In [16]: dr15.sort_values(ascending=False, by='importance').head(20)
```

```
Out[16]:
```

	feature	importance
0	age	0.190175
445	secs_elapsed	0.069931
411	first_browser_Chrome	0.029841
314	gender_MALE	0.028738
391	first_affiliate_tracked_untracked	0.028162
313	gender_FEMALE	0.027955
385	first_affiliate_tracked_linked	0.026072
315	gender_NotDisclosed	0.025224
416	first_browser_Firefox	0.022426
434	first_browser_Safari	0.020270
388	first_affiliate_tracked_omg	0.018865
317	signup_method_basic	0.015730
318	signup_method_facebook	0.015382
364	affiliate_channel_sem-brand	0.014952
461	device_type_Mac Desktop	0.013279
399	first_device_type_Mac Desktop	0.013251
375	affiliate_provider_google	0.012960
402	first_device_type_Windows Desktop	0.012765
465	device_type_Windows Desktop	0.012628
371	affiliate_provider_direct	0.012212

# Model Comparison

## ★ Pros:

- Decision Tree is very visual and can easily be interpretable.
- Random Forest was faster than Decision Tree.

## ★ Cons:

- Decision Tree is very slow. Needs lot of computing power.

# Future Scope

1. Lot more feature engineering (Season, Duration, Gender, Age etc.)
2. Removing features that are not important might help with the prediction and model run time
3. Random Forest might provide better result
4. Ensemble might help with accuracy further

For more info :

[https://github.com/psahu/GA\\_Project](https://github.com/psahu/GA_Project)

[https://github.com/subuone/sfdat22\\_work](https://github.com/subuone/sfdat22_work)

# Q & A

